


# Untuned model forecasts · 90 queries each

All six models complete · AI helps solve at least one Millennium Problem



Curves are medians of valid forecasts: 3 prompt variants × 30 distinct draws per model. All invalid outcomes are retained.  
Original prompts, BF16 runtime, sampling and output budgets are unchanged. Colors follow model release date.  
Extra 60 draws are baseline-only inference draws, not training seeds or new matched adapter comparisons.  
Models with incomplete collection are withheld. Low valid coverage limits interpretation; release-date prompts differ across models.