

Thinking on vs off: format sensitivity

POST-HOC SUPPLEMENT · One complete JSON code fence accepted · Frozen strict results preserved


All three comparisons complete · AI helps solve at least one Millennium Problem

—●— Thinking on —■— Thinking off

Qwen3 32B · 2025-04-29

Valid on: strict 90/90 → +fence 90/90


Valid off: strict 2/90 → +fence 90/90



Qwen3.5 27B · 2026-02-24

Valid on: strict 14/90 → +fence 14/90


Valid off: strict 89/90 → +fence 90/90



Qwen3.8 27B · 2026-08-14

Valid on: strict 4/90 → +fence 4/90

Valid off: strict 90/90 → +fence 90/90



Supplementary rule applied uniformly after observing formatting: complete single fenced JSON only; no prose, truncated or unclosed-thinking rescue.

Curves use the same valid matched draws; each mode schedules 3 prompt variants × 30 draws. These are generation draws, not training replicates.

Same weights, user prompts, RNG seeds and 4,096-token cap. The comparison includes thinking-token budget usage; it does not isolate internal reasoning.

Counts include completed mirrored batches only; live batches may contain more outputs. Original responses and the frozen strict analysis are preserved.